BRANDPILOT

Codex 48 — Round 2 Project Report

Evidence-linked AI brand governance before content goes live

Participant: Adit Dinesh Shinde

Stage: Round 2 — The Build

Status: MVP code package prepared; deployment and end-to-end verification must be completed before submission.

Do not claim cloud deployment or live AI verification until those steps have actually been completed.

1. Executive summary

BrandPilot is a pre-publication content review prototype. Users configure brand guidelines, preferred/avoided terminology, and approved facts; submit a draft; receive severity-ranked findings with explanations and suggested corrections; then record a human review decision and optionally save a revised draft for re-review.

2. Features implemented in the code package

- Streamlit interface with Review content, Review history, and About & limitations tabs.
- Editable brand guidelines, terminology pairs, and approved facts.
- Paste-in text review and .txt/.md upload.
- Deterministic terminology, risky claim, unsupported number, hype language, basic accessibility/inclusive-language, and readability checks.
- Finding cards include category, severity, quoted text, rationale, source/rule context, suggested correction, and confidence label.
- Human decision controls: Needs review, Changes requested, Approved by human, and reviewer comment.
- Editable revised draft and re-review flow.
- Local JSON review history.
- Optional Gemini API analysis path; deterministic checks remain available if AI is not configured or fails.

3. End-to-end workflow

Step	Behaviour
1. Configure	Set brand guidelines, terminology and approved facts in the sidebar.
2. Submit	Paste draft text or upload a .txt/.md file.
3. Analyse	Rule engine checks configured rules; optional Gemini adds contextual findings.
4. Explain	Inspect severity, flagged text, rationale, source/rule and suggested correction.
5. Decide	Human reviewer chooses a status and may add a comment.
6. Revise	Edit a revision and run the review again.
7. Inspect history	Review previous drafts, findings and decisions.

4. Technology and implementation

Python + Streamlit: UI and workflow. **Python regular expressions and rules:** terminology, risky claims, numeric claims, and readability heuristics. **Optional Google Gemini API via google-genai:** contextual review. **JSON file persistence:** local single-user review history. **Git/GitHub:** source control and distribution. **Streamlit Community Cloud:** recommended deployment target, pending actual deployment.

The AI layer is optional. When Gemini is configured, draft text and brand context are sent to the provider. Do not submit confidential material without authorisation and review of provider terms.

5. Changes from the Round 1 plan

The MVP narrows scope to text-first review. Visual brand comparison, complex document formats, enterprise-grade multi-user permissions, and CMS integrations are deferred. JSON persistence is used instead of a database for a faster single-user demo. AI is optional so judges can still use the central workflow without an API key.

6. Three major Round 3 improvements

1. **Visual brand and asset verification:** compare uploaded campaign images against approved logos/assets and flag likely misuse for human verification.

2. **Multi-user roles and secure persistence:** add authenticated creator/reviewer/admin roles, PostgreSQL storage, workspace isolation, and durable audit history.
3. **Document and workflow integrations:** add DOCX/PDF ingestion with source locations and a webhook/CMS draft handoff.

7. Testing checklist before submission

Verify that (1) the sample draft returns multiple findings; (2) changing terminology changes findings; (3) .txt upload works; (4) a human decision and comment save; (5) a revised draft can be re-reviewed; (6) history shows saved records; (7) the deployed link opens in a private/incognito window; and (8) AI failure does not break deterministic checks.

No user interviews, formal benchmark results, cloud deployment, or live AI performance are claimed here. Add only verified results after performing the tests.

8. Limitations and responsible use

This is a hackathon prototype. Approved-fact matching is heuristic and may miss or misflag claims. AI output can be inaccurate. Accessibility checks are not WCAG certification. Local JSON persistence is not suitable for concurrent production use. Human review is required before publication. Never commit API keys or confidential materials.

9. Links to add before submission

Repository URL: add the actual public GitHub repository after pushing the code. Working project URL: add the actual deployed Streamlit URL after deployment and testing. Do not submit placeholder URLs.