BRANDPILOT

Codex 48 — Round 2 Project Report

A text-first brand-governance MVP for reviewing draft content before publication.

Participant	Adit Dinesh Shinde
Project status	Deployed MVP; core workflow manually verified
Live application	brandpilot-adit-codex48.streamlit.app
Source repository	github.com/Adit24/BrandPilot

Verification note
This report distinguishes features present in the MVP from the specific behaviors manually verified on the deployed public app. It does not claim formal benchmarking, production readiness, WCAG certification, or verified Gemini-powered analysis.

Round 2 submission artifact • Updated to reflect the live deployment checks completed.

1. Executive summary

BrandPilot is a pre-publication content-review prototype that helps a reviewer identify potentially off-brand, exaggerated, or insufficiently supported claims in draft marketing content. A user can configure brand guidelines, preferred and avoided terminology, and approved facts; submit draft text; inspect severity-ranked findings and explanations; and record a human review decision with a comment.

The project is deployed on Streamlit Community Cloud and its public page was opened successfully in a private/incognito browser window. The rule-based review was exercised with a deliberately risky sample draft, and the human decision/history workflow was checked on the live deployment.

2. Implemented MVP features

Feature	Implementation
Brand workspace configuration	Editable brand name, brand guidelines, preferred/avoided terminology pairs, and approved facts.
Draft input	Paste text into the app and use the review workflow. The code package also supports .txt/.md uploads.
Rule-based content checks	Deterministic checks for risky absolute/promotional claims, hype language, terminology mismatches, and selected unsupported-number/readability/accessibility heuristics.
Explainable findings	Findings show severity/category, the flagged text, a rationale, and suggested correction/context where provided by the rule.
Human review decision	Reviewer can choose Needs review, Changes requested, or Approved by human, and optionally enter a comment.
Editable revision workflow	An edited draft can be saved as a new review for re-checking.
Review history	Saved review records display the draft, findings, decision, and reviewer comment.
Optional Gemini path	The code includes an optional Google Gemini analysis path when configured. It was not claimed as tested in this verification pass.

3. Core workflow

- **Configure:** Enter the brand voice, terminology rules, and approved facts.
- **Review:** Submit a draft for deterministic rule-based analysis.
- **Inspect:** Read findings and rationales, then decide whether changes are required.
- **Record:** Save a human decision and optional reviewer comment.
- **Track:** View saved review records and revisit the draft/findings/decision.

4. Manual verification completed

The following items were directly observed during testing of the public deployment. These are manual functional checks, not a formal benchmark or independent security audit.

Check	Result	Observed evidence
Public access	**PASS**	The deployed app opened in a private/incognito window without requiring the reviewer to sign in.
Risky-claim detection	**PASS**	Test draft: “Our revolutionary platform guarantees 100% risk-free results. We are the best in the world.” The app returned four findings: three high-severity promotional-claim findings (guarantees, 100% risk-free, best in the world) and one low-severity tone finding (revolutionary).
Human decision save	**PASS**	The app displayed the success message “Human decision saved.” after saving a reviewer decision/comment.
Correct decision in history	**PASS**	After selecting Changes requested and saving, the newest history entry displayed “Changes requested.”
Comment retention	**PASS**	The saved history entry displayed the comment: “Test: risky claims need evidence before approval.”
Review-history display	**PASS**	History displayed saved records, drafts, findings, decisions, and comments, including an earlier review with nine findings.
Persistence after reload	**PASS**	The user refreshed the browser and confirmed the saved Changes requested entry remained in history.
GitHub source upload	**PASS**	The push command succeeded for the main branch. The repository contains app.py, requirements.txt, README.md, .gitignore, and this report.

Not claimed as verified

- Gemini/API-based contextual analysis: optional in the code, but not confirmed as configured and tested.
- File upload, terminology changes, revised-draft re-review, and AI-failure fallback: implemented in the described MVP, but not included in the set of live checks confirmed above.
- No formal usability study, benchmark, accuracy measurement, load test, or security penetration test has been conducted.

5. Technical implementation

Python + Streamlit: interactive UI and workflow. **Python rules/regular expressions:** deterministic terminology and risky-claim checks plus lightweight heuristics. **Local JSON history:** saved single-user review records. **Git/GitHub:** source control. **Streamlit Community Cloud:** public demo deployment. **Optional Google Gemini API:** contextual analysis path when configured.

The deployed rule-based workflow does not require an API key. If Gemini is enabled, draft text and brand context may be sent to the configured provider; users should avoid entering confidential material unless authorized and comfortable with the provider's terms.

6. Changes from the Round 1 plan

- Scope was narrowed to a text-first MVP so the core review loop could be delivered and tested within the build window.
- A deterministic rule engine provides a usable baseline without requiring an external AI API.
- Review history uses local JSON persistence rather than a multi-user database.
- Visual asset comparison, enterprise authentication/permissions, and document/CMS integrations remain future work.

7. Three major Round 3 improvements

1. Visual brand and asset verification. Compare campaign images against approved logos/assets and flag possible misuse for human verification.

2. Secure multi-user review and audit trail. Add authentication, creator/reviewer/admin roles, workspace isolation, durable database storage, and auditable decision history.

3. Document ingestion and workflow integration. Support DOCX/PDF with source locations and add a controlled webhook/CMS handoff for draft review.

8. Limitations and responsible use

- This is a hackathon MVP, not a production-ready brand-governance system.
- Rule-based and approved-fact checks are heuristic; they may miss unsupported claims or flag legitimate content.
- Accessibility-related heuristics do not establish WCAG conformance or certification.
- Optional AI output can be inaccurate and must be reviewed by a person.
- Local JSON history is suitable for a single-user demo, not concurrent production use or a robust compliance audit.
- Human approval remains necessary before publication. Do not commit API keys, credentials, or confidential content.

9. Submission links

Live app: https://brandpilot-adit-codex48.streamlit.app/

Source repository: https://github.com/Adit24/BrandPilot

Submission statement
BrandPilot is a deployed text-first brand-governance MVP. The public access, rule-based risky-claim detection, human decision saving, review-history display, and persistence after page reload were manually verified. Other features are described as implemented in the code where applicable, but are not presented as independently tested unless stated above.